

# **Denoising Diffusion Probabilistic Models**

Jonathan Ho, Ajay Jain, and Pieter Abbeel

UC Berkeley

NeurIPS 2020

Presented by Minho Park

# Contribution

---

- Novel generative model which obtain sample quality similar to ProgressiveGAN.
  - A latent variable models training on a weighted variational bound.
- Connection between diffusion probabilistic models and denoising score matching with Langevin dynamics.
- Naturally admit a progressive lossy decompression scheme.

# Introduction

---

- A class of latent variable models inspired by considerations from nonequilibrium thermodynamics.




Figure 1: Generated samples on CelebA-HQ  $256 \times 256$  (left) and unconditional CIFAR10 (right)

# Background

---

- Notation
- Forward process:  $q(x_{0:T})$  - predefined Markov chain distribution by  $\mathcal{N}$ .
  - $q(x_0) := p_{data}(x_0)$ .
- Reverse process:  $p_\theta(x_{0:T})$  - Learnable distribution by  $\mathcal{N}$ .
- $\theta$ : Learnable parameters.




Figure 2: The directed graphical model considered in this work.

# Background

- Reverse process

$$p_{\theta}(x_{0:T}) := p(x_T) \prod_{t=1}^T p_{\theta}(x_{t-1}|x_t), \quad p_{\theta}(x_{t-1}|x_t) := \mathcal{N}(x_{t-1}; \mu_{\theta}(x_t, t), \Sigma_{\theta}(x_t, t))$$

- Forward process or diffusion process

$$q(x_{1:T}) := \prod_{t=1}^T q(x_t|x_{t-1}), \quad q(x_{t-1}|x_t) := \mathcal{N}(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$




Figure 2: The directed graphical model considered in this work.

# Background

---

- Optimizing the usual variational bound on negative log likelihood  $\mathbb{E}_{x_0 \sim q}[-\log p_\theta(x_0)]$ .

$$\begin{aligned} p_\theta(x) &= \int p_\theta(x_{0:T}) dx_{1:T} && \text{Marginal distribution} \\ &= \int p_\theta(x_{0:T}) \cdot \frac{q(x_{1:T}|x_0)}{q(x_{1:T}|x_0)} dx_{1:T} \\ &= \int p_\theta(x_T) \cdot \frac{\prod_{t=1}^T p_\theta(x_{t-1}|x_t)}{\prod_{t=1}^T q(x_t|x_{t-1})} \cdot q(x_{1:T}|x_0) dx_{1:T} \\ &= \int p_\theta(x_T) \cdot q(x_{1:T}|x_0) \cdot \prod_{t=1}^T \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} dx_{1:T} \end{aligned}$$

# Background

---

- Optimizing the usual variational bound on negative log likelihood  $\mathbb{E}_{x_0 \sim q}[-\log p_\theta(x_0)]$ .

$$\begin{aligned} & \mathbb{E}_{x_0 \sim q}[-\log p_\theta(x_0)] \\ &= \int -\log p_\theta(x_0) \cdot q(x_0) dx_0 \\ &= \int -\log \left( \int p_\theta(x_T) \cdot q(x_{1:T}|x_0) \cdot \prod_{t=1}^T \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} dx_{1:T} \right) \cdot q(x_0) dx_0 \\ &= \int -\log \left( \mathbb{E}_{q(x_{1:T}|x_0)} \left[ p_\theta(x_T) \cdot \prod_{t=1}^T \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \right] \right) \cdot q(x_0) dx_0 \\ &\leq \int -\mathbb{E}_{q(x_{1:T}|x_0)} \log \left( p_\theta(x_T) \cdot \prod_{t=1}^T \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \right) \cdot q(x_0) dx_0 \end{aligned}$$

Jensen's inequality

# Background

---

- Optimizing the usual variational bound on negative log likelihood  $\mathbb{E}_{x_0 \sim q}[-\log p_\theta(x_0)]$ .

$$\begin{aligned} & \int -\mathbb{E}_{q(x_{1:T}|x_0)} \log \left( p_\theta(x_T) \cdot \prod_{t=1}^T \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \right) \cdot q(x_0) dx_0 \\ &= \int -\log \left( p_\theta(x_T) \cdot \prod_{t=1}^T \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \right) \cdot \frac{q(x_0) \cdot q(x_{1:T}|x_0)}{q(x_{0:T})} dx_0 dx_{1:T} \\ &= \mathbb{E}_{x_{0:T} \sim q} \left[ -\log \left( p_\theta(x_T) \cdot \prod_{t=1}^T \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \right) \right] \\ &= \mathbb{E}_{x_{0:T} \sim q} \left[ -\log p_\theta(x_T) - \sum_{t=1}^T \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \right] \triangleq L \end{aligned}$$



# Background

---

$$\begin{aligned} L &= \mathbb{E}_{x_{0:T} \sim q} \left[ -\log p_{\theta}(x_T) - \sum_{t=1}^T \log \frac{p_{\theta}(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \right] \\ &= \mathbb{E}_{x_{0:T} \sim q} \left[ -\log p_{\theta}(x_T) - \sum_{t=2}^T \log \frac{p_{\theta}(x_{t-1}|x_t)}{q(x_t|x_{t-1})} - \log \frac{p_{\theta}(x_0|x_1)}{q(x_1|x_0)} \right] \\ &= \mathbb{E}_{x_{0:T} \sim q} \left[ -\log p_{\theta}(x_T) - \sum_{t=2}^T \log \frac{p_{\theta}(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} \cdot \frac{q(x_{t-1}|x_0)}{q(x_t|x_0)} - \log \frac{p_{\theta}(x_0|x_1)}{q(x_1|x_0)} \right] \\ &= \mathbb{E}_{x_{0:T} \sim q} \left[ -\log \frac{p_{\theta}(x_T)}{q(x_T|x_0)} - \sum_{t=2}^T \log \frac{p_{\theta}(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} - \log p_{\theta}(x_0|x_1) \right] \\ &= \mathbb{E}_{x_{0:T} \sim q} \left[ D_{KL}(q(x_T|x_0) \parallel p_{\theta}(x_T)) + \sum_{t=2}^T D_{KL}(p_{\theta}(x_{t-1}|x_t) \parallel q(x_{t-1}|x_t, x_0)) - \log p_{\theta}(x_0|x_1) \right] \end{aligned}$$

# Background

---

- Now the Evidence Lower BOund (ELBO) is tractable.

$$\mathbb{E}_{x_{0:T} \sim q} \left[ \underbrace{D_{KL}(q(x_T|x_0) \parallel p_\theta(x_T))}_{L_T} + \sum_{t=2}^T \underbrace{D_{KL}(p_\theta(x_{t-1}|x_t) \parallel q(x_{t-1}|x_t, x_0))}_{L_{t-1}} - \underbrace{\log p_\theta(x_0|x_1)}_{L_0} \right]$$

- with  $q(x_{t-1}|x_t, x_0) = \mathcal{N}(x_{t-1}; \tilde{\mu}_t(x_t, x_0), \tilde{\beta}_t I)$
- where  $\tilde{\mu}_t(x_t, x_0) := \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{(1-\bar{\alpha}_t)}x_0 + \frac{\sqrt{\bar{\alpha}_t}(1-\bar{\alpha}_{t-1})}{(1-\bar{\alpha}_t)}x_t$  and  $\tilde{\beta}_t := \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t}\beta_t$  (proof is started with Bayes' rule).

# Train DDPM

---

$$L_T = D_{KL}(q(x_T|x_0) \parallel p_\theta(x_T))$$

- $q(x_T|x_0)$  converges to standard Gaussian distribution.
- We assume  $p_\theta(x_T)$  as the standard Gaussian distribution.
- $\Rightarrow$  Do not have to train  $\theta$  in this term.

# Train DDPM

---

$$L_{t-1} = D_{KL}(p_\theta(x_{t-1}|x_t) \parallel q(x_{t-1}|x_t, x_0))$$

- $q(x_{t-1}|x_t, x_0) = \mathcal{N}\left(x_{t-1}; \underbrace{\frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{(1-\bar{\alpha}_t)}x_0 + \frac{\sqrt{\bar{\alpha}_t}(1-\bar{\alpha}_{t-1})}{(1-\bar{\alpha}_t)}x_t}_{\text{Estimate with } \mu_\theta(x_t, t)}, \underbrace{\frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t}\beta_t I}_{\text{All parameters are given.}}\right)$   
 $\therefore \Sigma_\theta = \frac{1-\bar{\alpha}_{t-1}}{1-\bar{\alpha}_t}\beta_t$
- Reparametrize with  $x_t(x_0, t) = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1-\bar{\alpha}_t}\epsilon$  for  $\epsilon \sim \mathcal{N}(0, I)$  and estimate  $\epsilon_\theta$  then minimize  $L_{t-1}$  can be same as


$$\mathbb{E}_{x_0, \epsilon} \left[ \frac{\beta_t^2}{2\sigma_t^2 \alpha_t (1 - \bar{\alpha}_t)} \|\epsilon - \epsilon_\theta(\sqrt{\bar{\alpha}_t}x_0 + \sqrt{1-\bar{\alpha}_t}\epsilon, t)\|^2 \right]$$

- which resembles denoising score matching over multiple noise scales indexed by  $t$ .

# Recap Score-base Generative Model

- We estimate  $p_{data}(x)$  to our dataset  $\{x_1, \dots, x_N\} \sim p_{data}(x)$ .

$$\arg \min_{\theta} \mathbb{E}_{p_{data}} \left[ \frac{1}{2} \|\nabla_x \log p_{data}(x) - s_{\theta}(x)\|_2^2 \right] \approx \arg \min_{\theta} \frac{1}{N} \sum_{i=1}^N \left[ \frac{1}{2} \|s_{\theta}(x)\|_2^2 + \text{trace}(\nabla_x s_{\theta}(x)) \right]$$



Estimated scores are only accurate in high density regions.

Estimated scores are accurate everywhere for the noise-perturbed data distribution due to reduced low data density regions.



# Train DDPM

---

$$L_0 = \mathbb{E}_{x_{0:T} \sim q} [-\log p_\theta(x_0|x_1)]$$

- We assume that image data consists of integers in  $\{0, 1, \dots, 255\}$  scaled linearly to  $[-1, 1]$ .
- Set the last term of the reverse process to an independent discrete decoder derived from the Gaussian  $\mathcal{N}(x_0; \mu_\theta(x_1, 1), \sigma_1^2 I)$ .

$$p_\theta(\mathbf{x}_0|\mathbf{x}_1) = \prod_{i=1}^D \int_{\delta_-(x_0^i)}^{\delta_+(x_0^i)} \mathcal{N}(x; \mu_\theta^i(\mathbf{x}_1, 1), \sigma_1^2) dx$$

$$\delta_+(x) = \begin{cases} \infty & \text{if } x = 1 \\ x + \frac{1}{255} & \text{if } x < 1 \end{cases} \quad \delta_-(x) = \begin{cases} -\infty & \text{if } x = -1 \\ x - \frac{1}{255} & \text{if } x > -1 \end{cases}$$

In practice, this term is optimized by MSE Loss.  
Furthermore, there is no independent decoder.

# Inference Algorithm

- Although training can be done by single step, **inference can not be done by single step.**




Figure 2: The directed graphical model considered in this work.

---

## Algorithm 1 Training

---

- 1: **repeat**
  - 2:  $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
  - 3:  $t \sim \text{Uniform}(\{1, \dots, T\})$
  - 4:  $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
  - 5: Take gradient descent step on  

$$\nabla_{\theta} \left\| \boldsymbol{\epsilon} - \boldsymbol{\epsilon}_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, t) \right\|^2$$
  - 6: **until** converged
- 

---

## Algorithm 2 Sampling

---

- 1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
  - 2: **for**  $t = T, \dots, 1$  **do**
  - 3:  $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$
  - 4:  $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$
  - 5: **end for**
  - 6: **return**  $\mathbf{x}_0$
-

# Simplified Training Objective

---

- Simplified objective aggregates  $L_{t-1}$  and  $L_0$ .
- Ignoring weight of  $t$  and discrete Gaussian distribution.

$$L_{simple}(\theta) := \mathbb{E}_{t, x_0, \epsilon} \left[ \left\| \epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon, t) \right\|^2 \right]$$

Table 2: Unconditional CIFAR10 reverse process parameterization and training objective ablation. Blank entries were unstable to train and generated poor samples with out-of-range scores.

Objective	IS	FID
<b><math>\tilde{\mu}</math> prediction (baseline)</b>		
$L$ , learned diagonal $\Sigma$	$7.28 \pm 0.10$	23.69
$L$ , fixed isotropic $\Sigma$	$8.06 \pm 0.09$	13.22
$\ \tilde{\mu} - \tilde{\mu}_{\theta}\ ^2$	–	–
<b><math>\epsilon</math> prediction (ours)</b>		
$L$ , learned diagonal $\Sigma$	–	–
$L$ , fixed isotropic $\Sigma$	$7.67 \pm 0.13$	13.51
$\ \tilde{\epsilon} - \epsilon_{\theta}\ ^2 (L_{simple})$	<b><math>9.46 \pm 0.11</math></b>	<b>3.17</b>



# Architecture

- U-Net backbone similar to an unmasked PixelCNN++.
- Parameter are shared across time, which is specified to the network using the Transformer sinusoidal positional embedding.
- Use self-attention at the  $16 \times 16$  feature map resolution.

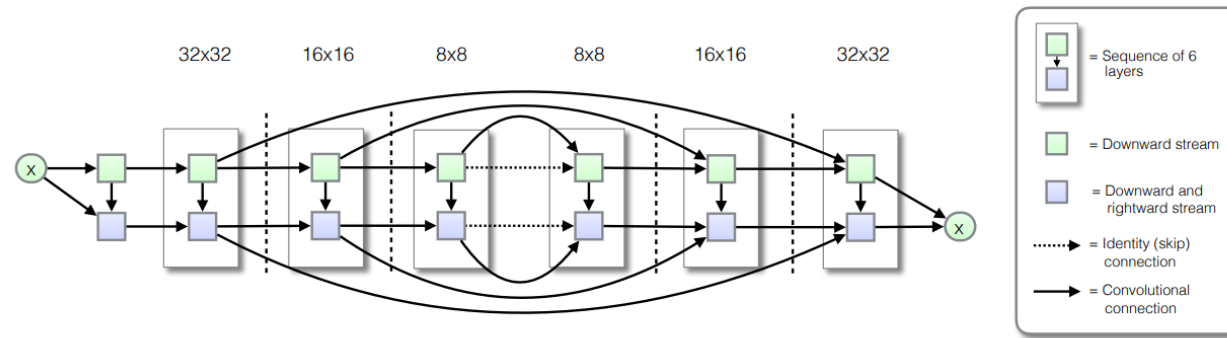


Figure 2: Like van den Oord et al. (2016c), our model follows a two-stream (downward, and downward+rightward) convolutional architecture with residual connections; however, there are two significant differences in connectivity. First, our architecture incorporates downsampling and upsampling, such that the inner parts of the network operate over larger spatial scale, increasing computational efficiency. Second, we employ long-range skip-connections, such that each  $k$ -th layer provides a direct input to the  $(K - k)$ -th layer, where  $K$  is the total number of layers in the network. The network is grouped into sequences of six layers, where most sequences are separated by downsampling or upsampling.

# Quantitative Results

Table 1: CIFAR10 results. NLL measured in bits/dim.

Model	IS	FID	NLL Test (Train)
<b>Conditional</b>			
EBM [11]	8.30	37.9	
JEM [17]	8.76	38.4	
BigGAN [3]	9.22	14.73	
StyleGAN2 + ADA (v1) [29]	<b>10.06</b>	<b>2.67</b>	
<b>Unconditional</b>			
Diffusion (original) [53]			$\leq 5.40$
Gated PixelCNN [59]	4.60	65.93	3.03 (2.90)
Sparse Transformer [7]			<b>2.80</b>
PixelIQN [43]	5.29	49.46	
EBM [11]	6.78	38.2	
NCSNv2 [56]		31.75	
NCSN [55]	$8.87 \pm 0.12$	25.32	
SNGAN [39]	$8.22 \pm 0.05$	21.7	
SNGAN-DDLS [4]	$9.09 \pm 0.10$	15.42	
StyleGAN2 + ADA (v1) [29]	<b><math>9.74 \pm 0.05</math></b>	3.26	
Ours ( $L$ , fixed isotropic $\Sigma$ )	$7.67 \pm 0.13$	13.51	$\leq 3.70$ (3.69)
<b>Ours (<math>L_{\text{simple}}</math>)</b>	$9.46 \pm 0.11$	<b>3.17</b>	$\leq 3.75$ (3.72)

# Qualitative Results



Figure 3: LSUN Church samples. FID=7.89



Figure 4: LSUN Bedroom samples. FID=4.90

# Progressive Coding

- Progressive lossy compression.
- Distortion (RMSE):  $\sqrt{\|x_0 - \hat{x}_0\|^2 / D}$  where  $\hat{x}_0 = \left(x_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_\theta(x_t)\right) / \sqrt{\bar{\alpha}_t}$

---

## Algorithm 3 Sending $\mathbf{x}_0$

---

```

1: Send  $\mathbf{x}_T \sim q(\mathbf{x}_T | \mathbf{x}_0)$  using  $p(\mathbf{x}_T)$ 
2: for  $t = T - 1, \dots, 2, 1$  do
3:   Send  $\mathbf{x}_t \sim q(\mathbf{x}_t | \mathbf{x}_{t+1}, \mathbf{x}_0)$  using  $p_\theta(\mathbf{x}_t | \mathbf{x}_{t+1})$ 
4: end for
5: Send  $\mathbf{x}_0$  using  $p_\theta(\mathbf{x}_0 | \mathbf{x}_1)$ 

```

---



---

## Algorithm 4 Receiving

---

```

1: Receive  $\mathbf{x}_T$  using  $p(\mathbf{x}_T)$ 
2: for  $t = T - 1, \dots, 1, 0$  do
3:   Receive  $\mathbf{x}_t$  using  $p_\theta(\mathbf{x}_t | \mathbf{x}_{t+1})$ 
4: end for
5: return  $\mathbf{x}_0$ 

```

---




Figure 5: Unconditional CIFAR10 test set rate-distortion vs. time. Distortion is measured in root mean squared error on a  $[0, 255]$  scale. See Table 4 for details.



# Progressive Coding

- Progressive generation.



Figure 6: Unconditional CIFAR10 progressive generation ( $\hat{\mathbf{x}}_0$  over time, from left to right). Extended samples and sample quality metrics over time in the appendix (Figs. 10 and 14).




Figure 7: When conditioned on the same latent, CelebA-HQ  $256 \times 256$  samples share high-level attributes. Bottom-right quadrants are  $\mathbf{x}_t$ , and other quadrants are samples from  $p_\theta(\mathbf{x}_0|\mathbf{x}_t)$ .

# Interpolation

- Large  $t$  results in coarser and more varied interpolations, with novel samples at  $t = 1000$ .




Illustration of interpolation with diffusion models

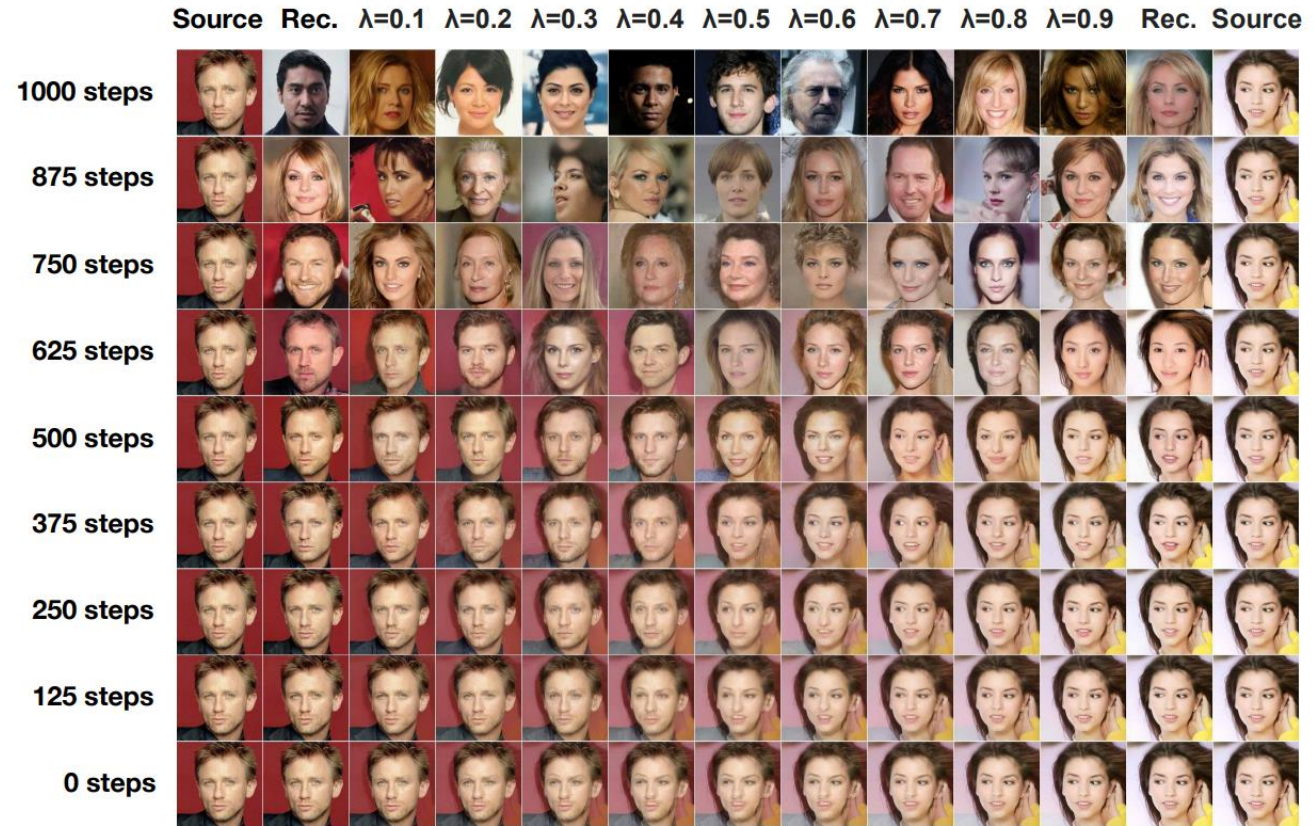


Figure 9: Coarse-to-fine interpolations that vary the number of diffusion steps prior to latent mixing.

# Contribution

---

- Novel generative model which obtain sample quality similar to ProgressiveGAN.
  - A latent variable models training on a weighted variational bound.
- Connection between diffusion probabilistic models and denoising score matching with Langevin dynamics.
- Naturally admit a progressive lossy decompression scheme.